1

Statistical methods for identifying hybrids and groups

Eric C. Anderson

Fisheries Ecology Division
Southwest Fisheries Science Center
110 Shaffer Road
Santa Cruz, CA 95060
USA

Introduction

Recently, statistical geneticists have developed a number of model-based methods that use genetic data to infer the population of origin of the gene copies within an individual. In this chapter we focus on three of these methods which are known by the software that implements them: *structure* (Pritchard et al. 2000), *NewHybrids* (Anderson and Thompson 2002), and *BayesAss+* (Wilson and Rannala 2003). These programs are increasingly used in animal conservation for population assignment, detection of hybridization, and estimation of recent migration rates. Unlike more generic statistical approaches (Bowcock et al. 1994; Roques et al. 2001), the three methods we review here are all based on an underlying probability model that is intended to mimic the inheritance of genes and the sampling of individuals. Such model-based inference has a number of advantages. First, it typically uses more of the information in the data than approaches that are not based explicitly on genetic models, and second, the variables appearing in genetically-based statistical models relate directly to genetic phenomena, so they are easily interpreted.

The statistical genetic models underlying *structure*, *NewHybrids*, and *BayesAss+* are simple and quite similar. The primary goal of this chapter is to describe these models with as few equations as possible. In lieu of mathematical equations we will explore the structure of these models in terms of simple, intuitive diagrams called *acyclic directed graphs*, DAGs, that show the relationship between variables in a model. This should allow users to better understand what the methods do, how they are similar, and the important ways in which they differ. Though the softwares implementing these techniques are user-friendly, they are certainly not “plug-and-play” methods. I hope that this presentation will allow users to understand the methods enough to ensure they get reasonable results and they can interpret them appropriately.

After discussing the three different models, we focus on practical issues. Because previous reviews (Pearse and Crandall 2004; Manel et al. 2005) have summarized when these various methods [and many related ones, (*e.g.*, Rannala and Mountain 1997; Dawson and Belkhir 2001; Piry et al. 2004; Corander et al. 2004)] are useful, and have offered many general guidelines for their use, this final section is devoted to the simple proposition that it is important to assess the results of these programs by comparison to simulated data that look like your own.

Conceptual models and graphical models

All statistical inference depends in some way on a probability model. This model may be completely specified in terms of the equations describing the statistical distributions involved; though if you simply want to understand the assumptions of the model, it is usually sufficient to understand the verbal description of the model. As a model gets more complex, however, it is helpful to have a visual roadmap as well as a verbal description. A DAG is such a roadmap, providing a diagram of the relationship between components in a model and a comparison of the structure of different models. We will illustrate our first DAG by considering the estimation of allele frequencies in a closed population.

Let us imagine that we are interested in estimating the frequency of alleles at a single locus in a lake population of fish. An obvious course of action would be to draw a sample of M fish from the lake, genotype them at the locus, and estimate the allele frequencies from the observed proportion of alleles in the sample. The conceptual model underlying this procedure is one in which each fish carries two gene copies drawn at random from a large pool of alleles whose proportions are the unknown allele frequencies in the lake.

The relationship between the allelic types in the fish we sample and the population allele frequencies is captured in the DAG of Figure 1.1(a). The circles (called nodes) in the graph represent the different variables in the model. The frequencies of the alleles are denoted by $\boldsymbol{\theta}$. The node associated with $\boldsymbol{\theta}$ is unshaded, representing the fact that the values of the population allele frequencies are unknown. The type of the allele is denoted by $Y_{i,1}$ for the first gene copy of the i^{th} sampled fish, and $Y_{i,2}$ for the second gene copy. The nodes associated with these variables are shaded black to denote that they are observed—*i.e.*, the fish are

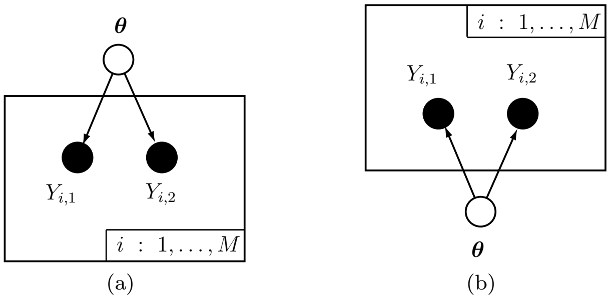

Fig. 1.1. DAGs describing the sampling to estimate allele frequency $\boldsymbol{\theta}$. a and b describe identical models.

genotyped. The allelic type of each gene copy is independent (under the assumption of Hardy-Weinberg equilibrium) and depends only on the frequency of alleles in the population. Hence, there are distinct arrows drawn from the node at $\boldsymbol{\theta}$ to the nodes for $Y_{i,1}$ and $Y_{i,2}$. The meaning of the arrow can be read as, for example, "the allelic type, $Y_{i,1}$, depends on $\boldsymbol{\theta}$." Finally, the two Y nodes are placed inside a box which is known as a *plate* (or, in this case, an M-plate). The legend at the lower right of the plate indicates that the variables within the plate are duplicated M times, over the subscript i. This shorthand expresses that M fish are sampled independently from the lake. The node for $\boldsymbol{\theta}$ is not included on the plate because each of the M fish is assumed sampled from the same population with the same allele frequencies. Figure 1.1(b) represents exactly the same model. This figure is included to emphasize that the spatial position of variables in the graph is unimportant, only the orientation of the arrows, and the connections they make, are relevant.

Recall that the original problem was to estimate the allele frequencies, $\boldsymbol{\theta}$, given the observed genotypes of a sample of fish. This problem is also apparent in the DAG because $\boldsymbol{\theta}$ is something we wish to know about, and yet it is unknown (as signified by its unshaded node). Generally, the problem of estimation can be interpreted in a DAG as the process of learning about variables or parameters with unshaded nodes given what is observed in the data (the shaded nodes).

As a final word on Figure 1.1, we should keep in mind that we would have obtained the same DAG if we were sampling any objects, two at a time, from a large population of objects. In fact, it is often easier to think of the sampling process as that of randomly drawing colored balls, two at a time, out of a large barrel. In this case, each ball is a gene copy, its color is its allelic type, and the barrel full of balls is the population of gene copies carried by fish in the lake. We have explored this example in detail because the "balls-in-barrels" conceptual model, and the DAG that goes with it, are basic building blocks for understanding more complex models. In the next section we use these building blocks to describe a class of models called *mixture models*.

Mixture models

The problem recently called "population assignment" in the molecular ecology literature is a special case of inference in a finite mixture model. In statistics, a finite mixture model is one in which the collection from which the sample is taken is a mixture of individuals from different populations. Such models were applied to the problem of population assignment and "genetic stock identification" as early as 1981 in the fisheries management literature (Milner et al. 1981). The programs *structure*, *NewHybrids*, and *BayesAss+* are all elaborations of the basic mixture model. In fact, the version of *structure* "without admixture" employs the same mixture model as an earlier method used to estimate proportions of Columbia river tributary salmon caught in a mixed stock fishery (Smouse et al. 1990).

This salmon-fishery mixture model arises from a scenario such as the following: K separate spawning populations of salmon, each with its own unknown allele frequencies, reproduce in different tributaries of a river. Fish from each population migrate through the same place (for example, the mouth of the river), where they are subject to a fishery. By sampling M fish in the fishery and genotyping them at L loci we hope to estimate the proportion of fish from each of the K tributaries that were at the fishery site when the sample was taken. We might also want to infer the population of origin of each of the sampled fish.

Figure 1.2(a) shows the DAG for the mixture model described above. This DAG is composed of a number of elements that look suspiciously like the DAG of Figure 1.1. Working our way through the graph, from top to bottom, we first have $\boldsymbol{\pi}$, which denotes the unknown proportions of fish from the K different populations at the fishery site. W_i is a variable that denotes the population of origin of the i^{th} fish. It can be thought of as a ball that is tied to the fin of the fish, with the color of the ball telling us where the fish comes from. Under this interpretation, each fish is sampled from the fishery as if it were a colored ball drawn

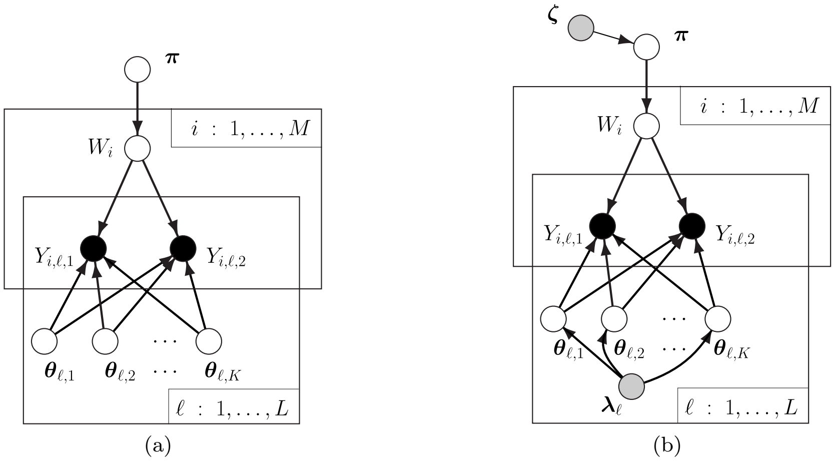

Fig. 1.2. DAGs for the mixture model. (a) represents the likelihood model, (b) includes nodes associated with the prior distributions for a Bayesian specification of the problem.

from a barrel in which the different colors of balls are in the unknown proportions $\boldsymbol{\pi}$. Of course, the node associated with W_i is unshaded because we don't know where the fish come from—that is what we would like to learn. The remainder of the graph looks complicated, but we can break it down as follows: the L-plate (the lower of the two plates, with the legend "$\ell : 1,\ldots,L$") signifies that for each fish, there are L loci genotyped, and that their allelic types are independent, given the fish's population of origin, W_i. $\boldsymbol{\theta}_{\ell,k}$ is the frequency of alleles at locus ℓ in population k (where k denotes any one of the K populations). $Y_{i,\ell,1}$ and $Y_{i,\ell,2}$ are the allelic types of the gene copies carried by fish i at locus ℓ. As the arrows in the graph show, the allelic types of these gene copies depend *both* on W_i and on the allele frequencies in the different populations. The nature of this dependence is straightforward: the two allelic types $Y_{i,\ell,1}$ and $Y_{i,\ell,2}$ are drawn from the population that the fish is from—which is denoted by W_i.

The whole sampling model can be summarized by thinking about generating a sample from it. The steps in doing so would be: 1) draw a ball from a barrel with frequencies $\boldsymbol{\pi}$. 2) The color of the ball tells you which population to sample a fish from. 3) To generate the genotype for that fish you draw two balls, representing the genes in that fish, from each of L barrels. Each barrel represents the alleles in the population at one of the L loci.

As before, the inference problems that can be tackled with this model can be seen in the DAG. The exercise of population assignment (Paetkau et al. 1995; Rannala and Mountain 1997) is merely that of inferring the values of the W_i variables. On the other hand, estimating the proportion of fish from different populations is just the process of inferring the value of $\boldsymbol{\pi}$. Finally, if desired, one could also pursue inference of the allele frequencies in the populations. These are all inference problems that are just different uses of the same underlying model. In actual practice, these different inference problems are typically tackled at the same time, but it is still useful to view them as separate inference problems.

Many times, individuals known to be from particular populations may be sampled. Such individuals constitute what are called *learning samples* or *training samples.* These would be represented in the DAG simply as individuals for whom the node associated with W_i was shaded. Inference then proceeds much as before—unknown quantities of interest are estimated given the observed data, which in this case includes the W_i's of the individuals in the learning samples. Though with multiple loci mixture inference may be possible without learning samples, if there are many populations contributing to the mixture then accurate inference may be impossible without learning samples.

Bayesian Inference

structure, *NewHybrids*, and *BayesAss+* all use the Bayesian paradigm for inferring quantities of interest. This means that estimation is conducted by summarizing the posterior distribution of quantities of interest. The posterior distribution of an unknown variable is just its probability distribution conditional on the observed data. Computing the posterior distribution can be difficult, and, indeed, in *structure*, *NewHybrids*, and *BayesAss+* it is approximated using Markov chain Monte Carlo. However, the fact that the inference is done in a Bayesian manner does not substantially alter the structure of the underlying models. This is illustrated in Figure 1.2(b), which shows the DAG for a Bayesian specification of the mixture model of Figure 1.2(a). It is apparent that the "heart" of the model is unchanged. In fact, the only modification is the addition of prior distributions parametrized by $\boldsymbol{\zeta}$ for $\boldsymbol{\pi}$ and $\boldsymbol{\lambda}_\ell$ for the $\boldsymbol{\theta}_\ell$'s. The nodes for $\boldsymbol{\zeta}$ and $\boldsymbol{\lambda}_\ell$ are shaded grey to denote that values of those parameters are assumed rather than observed. Prior distributions are necessary for Bayesian inference. Usually the parameters of the prior distribution are chosen to reflect prior knowledge—or in many cases, ignorance—about the associated variables.

A Survey Of Methods

Having established the language of graphical models, we are now in position to quickly survey the models used in *structure*, *NewHybrids*, and *BayesAss+*.

The structure *model without admixture*

As indicated above, the *structure* model without admixture is identical to the model shown in Figure 1.2(b), and the details of that model have already been described. It assumes that all individuals descend exclusively from one of K populations, where K can be set by the user. In other words, there is no facility in this model for explicitly dealing with hybrids or admixed individuals. Therefore, the method should be used with collections of organisms that are believed to be non-interbreeding. The data required are the multilocus genotypes of the individuals in a sample. The individuals may belong to "cryptic" subpopulations. That is, it is not necessary to have prior knowledge of separate groups—the program will automatically infer K subpopulations; however, the inclusion of learning samples can be helpful in resolving groups, especially if K is large, or genetic differentiation between populations is limited. The program computes the posterior probability that each individual belongs to each of the K subpopulations, and, in the process it also estimates the allele frequencies in the K separate subpopulations.

It is worth noting that when *structure* uses the model with no admixture, it assumes that the proportion of individuals from each subpopulation is equal (each subpopulation contributes a proportion $1/K$ to the mixture). This feature will cause *structure* to overestimate the true posterior probability of group membership for individuals from subpopulations that are rare in the mixture. If this is a concern, then it may be preferable to use the program *BAYES* (Pella and Masuda 2001) which was developed for analyzing large mixtures of salmon.

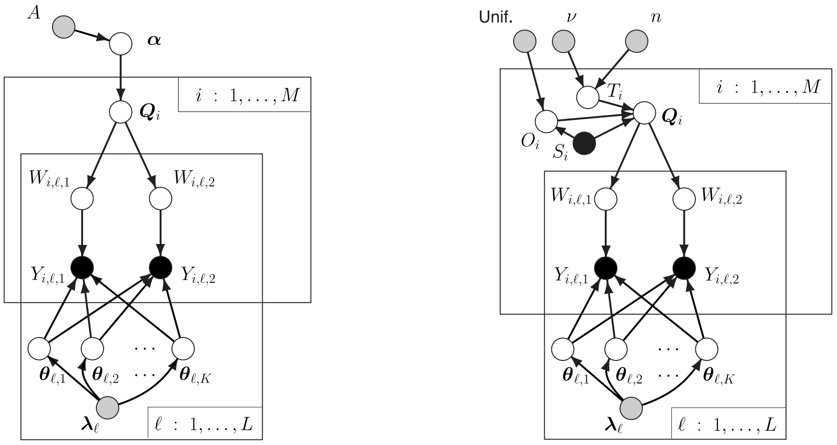

(a) (b)

Fig. 1.3. (a) The *structure* model with admixture. (b) The *structure* model with admixture *and* prior population information.

The **structure** *model with admixture*

This model provides a flexible way of accommodating individuals of mixed ancestry. No longer must each individual be purely descended from one of the K subpopulations. Rather, each gene copy within an individual may come from a different one of the K subpopulations. The subpopulations of origin of the two gene copies at locus ℓ in the i^{th} individual are indicated by the unobserved variables $W_{i,\ell,1}$ and $W_{i,\ell,2}$, and the expected proportion of ancestry of the i^{th} individual from each of the K subpopulations is a variable to be inferred, denoted by $\boldsymbol{Q}_i$. The DAG for this model appears in Figure 1.3(a). α is a parameter that determines whether individuals tend to be mostly admixed (high values of α) or mostly purebred (low values of α). It is a value that can be assumed, or inferred. If it is inferred, its prior distribution is assumed to be uniform on the interval $(0, A]$.

We can use the DAG to follow how we would generate data under the model, given α and the allele frequencies: 1) Conditional on α we would simulate a different, random $\boldsymbol{Q}_i$ for each individual i in the sample. $\boldsymbol{Q}_i$ can be thought of as the proportion of balls of K different colors filling a "Q-barrel" associated with individual i. 2) For each locus, we would draw two balls from individual i's Q-barrel. The colors of the balls drawn tell us which of the K different subpopulations the two gene copies at each locus came from. 3) The allelic type of each gene copy would then be drawn from the allele frequencies in the gene copy's subpopulation of origin.

This is a flexible and general model. It applies generically to many different scenarios: estimating the hybrid index (*i.e.*, $\boldsymbol{Q}_i$) of individuals in hybrid zones, detecting recent gene flow between populations, and elucidating population structure (cryptic or otherwise). It also provides a facility for estimating the number of subpopulations in a structured population, without prior knowledge about population boundaries.

The data required are the multilocus genotypes of sampled individuals. Learning samples are not required, so it is possible to identify cryptic genetic population structure in a sample of individuals from a single location. However, the capacity to detect cryptic structure declines as the degree of admixture of the individuals in the sample increases (Falush et al. 2003). In other words, if most individuals in the sample are highly-admixed members of a hybrid swarm, it will be more difficult to correctly infer the nature of the population structure than if some of the individuals in the sample retain the genotypes of pure subspecies, and others are admixed.

The **structure** *model with admixture and prior population information*

A variant available with *structure* is the model with "prior population information" in which genotyped individuals have been sampled from K known, separate subpopulations. This model is used to identify individuals in each sample that are migrants from other subpopulations or that have recent immigrant ancestry. In this case, it is necessary to have prior knowledge that there are distinct subpopulations, and that K of them have been sampled. A subpopulation is typically comprised of individuals living in a particular locality; however, the definition of "subpopulation" is flexible. For example, one might be able to define K subpopulations on the basis of distinct morphological traits possessed by different species or subspecies.

Figure 1.4(a) is a schematic of the population model in the case of $K = 3$ subpopulations of cats. The subpopulations are distinct, but there is migration between them. Immigration is assumed to be symmetrical and equal between all subpopulations. The model specifies that each individual has a probability $1 - \nu$ of being descended purely from ancestors belonging to the subpopulation from which it was sampled.

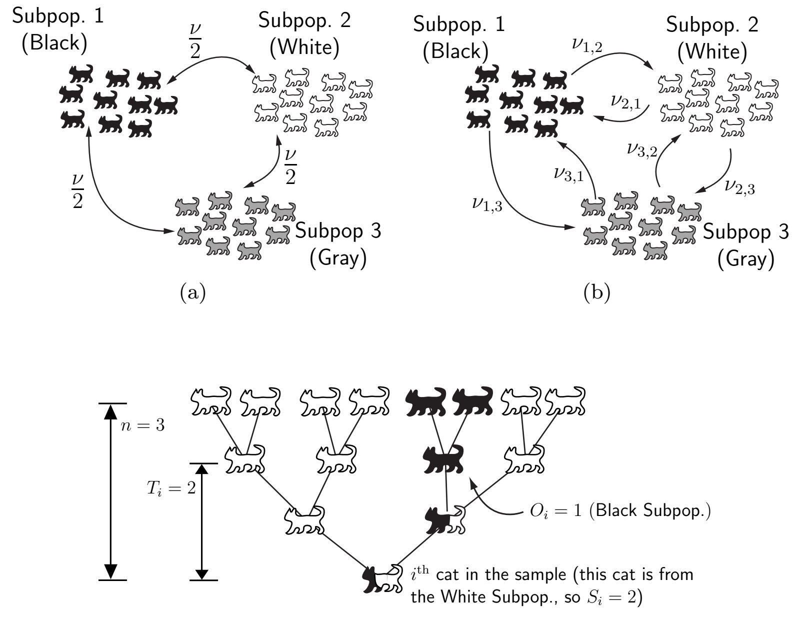

(a) (b)

(c)

Fig. 1.4. (a) A schematic of the *structure* model with prior population information assuming three subpopulations of cats. ν is the fraction of individuals in any subpopulation having a single immigrant ancestor in the last n generations from any of the other subpopulations. The other subpopulations are assumed to contribute migrants at the same rate so the probabiilty that an individual has an ancestor from a specific subpopulation is $\nu/(K-1)$, which in this case is $\nu/2$ because there are $K = 3$ subpopulations. (b) The migration model in *BayesAss+*. $\nu_{j,k}$ is the fraction of individuals in subpopulation k having immigrant ancestors from subpopulation j in the last n generations. (c) Notation relating to migrants and their descendants. S_i is the location where cat i was sampled. T_i is the number of generations back in time that i had a single migrant ancestor. O_i is the origin of that migrant. n is the total number of generations in the past during which it is assumed an individual might have a migrant ancestor. The cat shown at the bottom of the pedigree was sampled from the White Subpopulation ($S_i = 2$) and it has a single migrant ancestor from the Black Subpopulation ($O_i = 1$) two generations ago ($T_i = 2$). Correspondingly, it is expected to have $\frac{1}{4}$ of its ancestry from the Black Subpopulation, $\frac{3}{4}$ from the White Subopulation, and no ancestry from the Gray Subpopulation. *i.e.*, $\boldsymbol{Q}_i = (\frac{1}{4}, \frac{3}{4}, 0)$

With probability ν, however, an individual has immigrant ancestry. If ν is unknown (as it usually is) then it must be assumed.

If individual i has immigrant ancestry, then it is assumed that only one ancestor in the last n generations was a migrant, and that this migrant ancestor arrived from subpopulation O_i in the T_i^{th} generation before sampling. If $T_i = 0$ then the sampled individual i is itself the migrant; if $T_i = 1$ then one of individual i's two parents was a migrant; if $T_i = 2$ then one of i's four grandparents was a migrant, and so forth (Figure 1.4(c)). O_i and T_i are unknown. We will let S_i denote the subpopulation from which the i^{th} individual was sampled; S_i is an observed variable.

The DAG in Figure 1.3(b) shows that this model with prior population information is identical to the original *structure* model, except for the parts "upstream" from the $\boldsymbol{Q}_i$ node. In effect, the model with admixture and prior population information just establishes a new, and more easily interpreted, prior probability distribution for $\boldsymbol{Q}_i$ that ultimately depends on ν and n. The arrows in the DAG appear as they do because 1) the parameters ν and n determine the probability that an individual has a migrant ancestor at time T_i; 2) if individual i has a migrant ancestor, then the origin of that migrant depends on S_i because the migrant must have come from somewhere *other* than S_i; and finally 3) given T_i, S_i, and O_i, the value Q_i is determined (Figure 1.4(c)).

This specialized model is tailored to provide more power than the generic *structure* model for detecting individuals with recent immigrant ancestry. The data required are the multilocus genotypes of the sampled individuals *and* knowledge of the subpopulation each individual was sampled from. Being more specialized, this model also makes more assumptions. Specifically, it is assumed that migration occurs infrequently at a known rate, and that migration occurs at the same rate from and into all subpopulations. This is a model for detecting migrants; not for detecting non-migrants. It is important to note that the way the model is set up, if there were no genetic data, the posterior probability that a individual is *not* a migrant is $1 - \nu$. Therefore, if you choose ν to be .01 and run *structure* to discover that the posterior probability that each individual in your sample is a non-migrant is .99, you *must not* infer that this is telling you anything about the power of your genetic data to distinguish the subpopulations—you would have gotten the same result even if you had no genetic data.

Looking at the DAG of Figure 1.3(b), it might not be immediately clear how the genetic data, $Y_{i,\ell,1}$ and $Y_{i,\ell,2}$, will influence the posterior distribution of $\boldsymbol{Q}_i$—after all, there are no arrows from $Y_{i,\ell,1}$ or $Y_{i,\ell,2}$ to $\boldsymbol{Q}_i$, so how can $\boldsymbol{Q}_i$ depend on $Y_{i,\ell,1}$ or $Y_{i,\ell,2}$? The answer is that, even though it is natural in the formulation of a probability model to speak of one variable depending on another—for example, the color of a ball drawn from a barrel depends on the frequency of different-colored balls in the barrel—the influence between variables runs in both directions along the arrow. This is, in fact, why it is possible to do inference: if most of the balls you draw from a barrel are orange, then you may infer that there is a high frequency of orange balls in the barrel in the barrel. In other words, the observed data influence your belief about unobserved variables. In the case of the *structure* model of Figure 1.3(b), knowing the allelic type $Y_{i,\ell,1}$ gives you some information about where that gene copy came from ($W_{i,\ell,1}$) if you have some idea about the allele frequencies. Information about $W_{i,\ell,1}$, in turn, influences your belief about $\boldsymbol{Q}_i$ which, in turn, influences your belief about T_i and O_i which are variables that describe whether an individual is a migrant or not. In other words, during the inference process information obtained from observed variables flows throughout the graph to influence one's belief about *all* the unobserved variables and parameters. A corollary is that with no data, the posterior distribution of variables or parameters will merely be their prior distribution, *i.e.*, with no genetic data, the posterior probability that an individual is a migrant is merely its prior probability, ν.

There are two important limitations of the *structure* model with admixture and prior population information. The first is that it does not account for the fact that descendants of migrants will inherit genes in predictable *patterns* (not just in predictable *proportions*) from the different subpopulations (more details appear in the following section). The second limitation is the requirement that the migration rate ν must be known, or assumed. It would be preferable to allow the estimation of ν from the data.

The **NewHybrids** *model*

NewHybrids is designed to identify individuals that are recent hybrids between two species or populations. It can distinguish between genealogical classes like F1, F2, and backcrosses in a way that *structure* cannot because the *NewHybrids* model takes account of the predictable *patterns* of gene inheritance in hybrids, while the *structure* model does not. The simplest example occurs in comparing F1 hybrids (the offspring of parents from different populations or species) with F2 hybrids (the offspring of two parents who are themselves F1 hybrids). F1 hybrids will have *exactly* one gene copy from one population and one gene copy from the other population *at every locus*. An F2 individual will also have, on

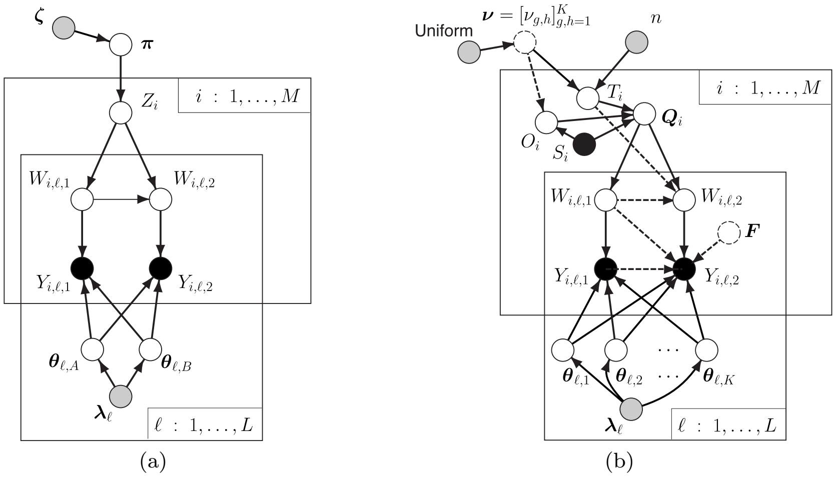

Fig. 1.5. (a) The NewHybrids model. (b) The model used in *BayesAss+*. The additions to the model that make it different from *structure* with admixture and prior population information are depicted with dashed lines.

average, half of its gene copies from one population and half from the other; however, only in half of its loci, on average, will there be exactly one gene copy from each population. The model in *structure* is not able to detect differences between F1's and F2's because it models admixture strictly in terms of $\boldsymbol{Q}_i$, which is the proportion of gene copies an individual will have, *on average*, from different subpopulations.

The DAG for the *NewHybrids* model (Figure 1.5(a)), shows that it is a mixture model. In this case, however, the different components of the mixture are different genealogical classes, rather than simply different populations. $\boldsymbol{\pi}$ denotes the proportions of individuals of different genealogical classes present where the sample is drawn, and Z_i is an unobserved variable that denotes the genealogical class of individual i. There are only two different species or populations (A and B) that an individual's genes may come from.

The model can be described by imagining simulating data from it, conditional on $\boldsymbol{\pi}$ and the allele frequencies. For the i^{th} individual: 1) a colored ball is drawn from a barrel with balls in the proportions of $\boldsymbol{\pi}$. The color of the ball gives the genealogical class (Z_i) of the individual; 2) given the genealogical class, the population of origin of the individual's first gene copy ($W_{i,\ell,1}$) is drawn from a barrel much like the Q-barrel described before. 3) The origin of the second gene copy ($W_{i,\ell,2}$) is drawn from a distribution that depends not only on the genealogical class, but also on the origin of the first gene copy. For example, if the genealogical class is F1, and the first gene copy came from population A, then the second gene copy must come from population B. 4) The allelic type of each gene copy is drawn from the allele frequencies in their respective populations of origin.

Visible in the DAG are the inference problems that can be tackled with *NewHybrids*. $\boldsymbol{\pi}$ can be estimated, and the genealogical class of each individual in the sample can be inferred. Also, the allele frequencies in populations A and B may be estimated.

The number of genealogical classes used in *NewHybrids* can be determined by the user. The default is 6: two pure species categories, F1, F2, A-backcross, and B-backcross categories. A considerable amount of genetic data is required to distinguish genealogical classes, even with as few as 6 classes (Vähä and Primmer 2006). It is even more difficult to resolve other genealogical classes like second- or third-generation backcrosses. Hence, *NewHybrids* is particularly appropriate for the study of hybrid zones in which hybridization has started to occur only recently, or in which the degree of introgression and backcrossing is limited due to selection against hybrids. It is worth pointing out that if only the two pure categories (Pure A and Pure B) are used, the *NewHybrids* model reduces to the standard mixture model of Figure 1.2(b) with $K = 2$.

The data required are the multilocus genotypes of the sampled individuals. Learning samples are not necessary, but they may be included. It is not necessary to have prior information about subpopulations or species.

The BayesAss+ *model*

The model in *BayesAss+* is a natural extension of the *structure* model with admixture and prior population information. Figure 1.4(b) gives a schematic of the migration model. Importantly, migration rates are not constrained to be the same between all pairs of populations. Further, with *BayesAss+* it is not necessary to assume a value of the migration rate. Rather, *BayesAss+* endeavors to estimate the (possibly non-symmetrical) rates of migration between all subpopulations. The other two advances over *structure* are the correct modeling of patterns of gene inheritance and the inclusion of an inbreeding parameter $\boldsymbol{F} = (F_1, \ldots, F_K)$ that tries to account for possible departures from Hardy-Weinberg equilibrium within each subpopulation.

Comparing the DAG for the *BayesAss+* model (Figure 1.5(b)) to that of *structure* with admixture and prior population information (Figure 1.3(b)) shows that the two are similar, differing only in a few variables, and a few extra arrows. Proceeding from top to bottom in the DAG, we first see that ν has been replaced with a matrix $\boldsymbol{\nu}$ of individual migration rates between the populations (Figure 1.4(b)). There is a new arrow connecting $\boldsymbol{\nu}$ to O_i because, since immigration rates are no longer symmetrical and equal, the origin of immigrants depends both on their destination S_i *and* on the migration matrix $\boldsymbol{\nu}$. The two new arrows, from T_i and $W_{i,\ell,1}$ to $W_{i,\ell,2}$ are there as a consequence of the fact that *BayesAss+* models the inheritance of genes from migrants in the same way that *NewHybrids* does genes in F1's and backcrosses. Finally, the arrows from $W_{i,\ell,1}$, $Y_{i,\ell,1}$, and $\boldsymbol{F}$ to $Y_{i,\ell,2}$ describe the interdependence of those variables induced by the possibility of inbreeding (departures from Hardy-Weinberg equilibrium). In words, the type of the second gene copy at a locus is no longer independent of the type of the first gene copy even if they both originate from the same population.

The primary goal of inference using this model is the estimation of the migration matrix. The data requirements for *BayesAss+* are the same as they are for *structure* with admixture and prior population information—it requires multilocus genotypes sampled from K distinct subpopulations. The model provides a more faithful representation of the data than does *structure* and it is appropriate for estimating recent migration between populations that are well-differentiated, genetically. However, it is apparent that if the populations are not greatly differentiated, then it may be difficult to estimate the migration rates between them. This could lead to misleading results if attempting to estimate migration rates between demes of a recently fragmented population. The various demes will be similar genetically due to recent common ancestry, and this might lead to inflated estimates of migration rates, even if no migration is presently occurring due to the recent fragmentation. Similarly, users should be suspicious of non-migration rates close to $\frac{2}{3}$ as this is the minimum allowed by the program and may indicate that populations are not genetically differentiated to the level required to get reliable results.

Practical Issues

Quite reasonably, an entirely separate chapter could be written dealing with practical issues involved in running the programs described here;

issues from "How large should my samples be?" and "How many loci should I use?" to issues like "Why do I get different results in *NewHybrids* using different priors for the allele frequencies?" and "Can I trust the results from these programs?" While some recent simulation studies (Evanno et al. 2005; Vähä and Primmer 2006) have addressed these sorts of questions, and have provided some general answers, the behavior of these methods is affected by many different features of the data, including the genetic differentiation between the populations, the number of alleles at each locus, the degree of admixture, *etc.* It is unlikely that any simulations that have been done will correspond well to all such features in your own data set. Furthermore, you may have different questions in mind than the ones that were addressed in any particular simulation study. In such cases, it is valuable to compare your results to the results obtained by analyzing data simulated to look like your own data set under different hypotheses of interest.

An excellent example of this type of effort appears in an analysis of structure in cod populations in the seas around Denmark (Nielsen et al. 2003). The authors were interested in whether the patterns of genotypes they observed in a contact zone were concordant with mechanical mixing of pure members of two populations, or with a zone of admixture between two populations. This is not a question that *structure* automatically addresses, so the two different scenarios were simulated with a program called *HybridLab* (see Nielsen et al. 2003 for details of the program) using allele frequencies from the two different pure populations. The results from the simulated admixture scenario were more similar to the results from the real data than were the results from the simulated mechanical mixing scenario, providing evidence that admixture between the populations may be occurring.

Simulating multilocus genotype data from specific allele frequencies is not a difficult task, but is not a standard feature in many genetic simulation programs. In addition to *HybridLab*, the program *spip* (Anderson and Dunham 2005) simulates multilocus genotypes from specified allele frequencies, and the program *simdata_nh* (available from santacruz.nmfs.noaa.gov/staff/eric_anderson/) simulates genotypes of individuals of different genealogical classes under the *NewHybrids* model. These programs can be used to test the inferences from the three programs *structure*, *NewHybrids*, and *BayesAss+*.

The methods reviewed in this chapter are complex enough that it is difficult (even for the authors of the programs) to make specific predictions about how these methods will behave when confronted with specific data sets. For this reason, the most important practical advice I can give is that it is incumbent upon the careful user of these programs to simulate data that is similar to their own and then analyze it with the program they are using. In order to gain insight about the results of these programs, there really is no substitute for comparing your results to the results achieved using simulated data that look like your own, *but in which you know the truth* (*i.e.*, you know which individuals are F1's, and F2's, or which ones are migrants).

Acknowledgments

I thank Devon Pearse, Ingrid Spies, and Greg Wilson for their insightful comments on a draft of this chapter, and Robin Waples and Kristen Ruegg for helpful discussions on the topic.

References

Anderson, E. C. and K. K. Dunham (2005). spip 1.0: a program for simulating pedigrees and genetic data in age-structured populations. *Molecular Ecology Notes 5*, 459–461.

Anderson, E. C. and E. A. Thompson (2002). A model-based method for identifying species hybrids using multilocus genetic data. *Genetics 160*, 1217–1229.

Bowcock, A. M., A. Ruiz-Linares, J. Tomfohrde, E. Minch, J. R. Kidd, and L. L. Cavalli-Sforza (1994). High resolution of human evolutionary trees with polymorphic microsatellites. *Nature 368*, 455–457.

Corander, J., P. Waldmann, P. Marttinen, and M. J. Sillanpaa (2004). BAPS 2: enhanced possibilities for the analysis of genetic population structure. *Bioinformatics 20*, 2363–2369.

Dawson, K. J. and K. Belkhir (2001). A Bayesian approach to the identification of panmictic populations and the assignment of individuals. *Genetical Research 78*, 59–77.

Evanno, G., S. Regnaut, and J. Goudet (2005). Detecting the number of clusters of individuals using the software STRUCTURE: a simulation study. *Molecular Ecology 14*, 2611–2620.

Falush, D., M. Stephens, and J. K. Pritchard (2003). Inference of population structure using multilocus genotype data: linked loci and correlated allele frequencies. *Genetics 164*, 1567–1587.

Manel, S., O. E. Gaggiotti, and R. S. Waples (2005). Assignment methods: matching biological questions with appropriate techniques. *Trends in Ecology & Evolution 20*, 136–142.

Milner, G. B., D. J. Teel, F. M. Utter, and C. L. Burley (1981). Columbia River stock identification study: validation of method. Technical report, NOAA, Northwest and Alaska Fisheries Center, Seattle, Washington.

Nielsen, E. E., M. M. Hansen, D. E. Ruzzante, D. Meldrup, and P. Gronkjaer (2003). Evidence of a hybrid-zone in Atlantic cod (*Gadus morhua*) in the

Baltic and the Danish Belt Sea revealed by individual admixture analysis. *Molecular Ecology 12*, 1497–1508.

Paetkau, D., W. Calvert, I. Stirling, and C. Strobeck (1995). Microsatellite analysis of population structure in Canadian polar bears. *Molecular Ecology 4*, 347–354.

Pearse, D. E. and K. A. Crandall (2004). Beyond F-ST: Analysis of population genetic data for conservation. *Conservation Genetics 5*, 585–602.

Pella, J. and M. Masuda (2001). Bayesian methods for analysis of stock mixtures from genetic characters. *Fisheries Bulletin Seattle 99*, 151–167.

Piry, S., A. Alapetite, J.-M. Cornuet, D. Paetkau, L. Baudouin, and A. Estoup (2004). *GeneClass2*: A software for genetic assignment and first-generation migrant detection. *Journal of Heredity 95*, 536–539.

Pritchard, J. K., M. Stephens, and P. Donnelly (2000). Inference of population structure using multilocus genotype data. *Genetics 155*, 945–959.

Rannala, B. and J. L. Mountain (1997). Detecting immigration by using multilocus genotypes. *Proceedings of the National Academy of Sciences, USA 94*, 9197–9201.

Roques, S., J. M. Sevigny, and L. Bernatchez (2001). Evidence for broadscale introgressive hybridization between two redfish (genus *Sebastes*) in the North-west Atlantic: a rare marine example. *Molecular Ecology 10*, 149–165.

Smouse, P. E., R. S. Waples, and J. A. Tworek (1990). A genetic mixture analysis for use with incomplete source population data. *Canadian Journal of Fisheries and Aquatic Science 47*, 620–634.

Vähä, J.-P. and C. R. Primmer (2006). Efficiency of model-based methods for detecting hybrid individuals under different hybridization scenarios and with different numbers of loci. *Molecular Ecology 15*, 63–72.

Wilson, G. A. and B. Rannala (2003). Bayesian inference of recent migration rates using multilocus genotypes. *Genetics 163*, 1177–1191.